

WORKING PAPER**The AI Risk Economy***Why Insurance Cannot Price What Governance Cannot Prove*

Basil C. Puglisi, MPA

A Human-AI Collaboration

basilpuglisi.com

Working Paper v5.0 · May 12, 2026

AI ethics, responsible AI, and AI governance are not synonyms, even though the industry has been treating them as if they were. In January 2026, the insurance market made the difference between them legally and financially visible. The companies that have actual AI governance will be able to defend a claim when something goes wrong. The companies that only have ethics statements and responsibility policies will be exposed. The premium difference between those two outcomes is about to show up in renewal cycles across 2026 and 2027.

The three terms describe different layers of a single problem. Ethical AI is the boundary, the conversation about what AI should and should not be allowed to do, where the line sits between acceptable risk and unacceptable harm. It produces principles and white papers, but not the kind of accountability that survives a renewal audit. Responsible AI is the engineering layer: testing, documentation, bias detection, hallucination monitoring, and human oversight built into the code. It produces an audit trail of how a system was built, not a named human accountable for any decision the system later made. AI Governance fills the gap the first two leave open. It is a named human with binding authority at a defined decision point, with the information needed to decide, with a record of what was decided and why. It is the layer that an insurance adjuster, a regulator, a journalist, or a plaintiff's lawyer can audit when something goes wrong.

Companies have been collapsing all three into a single phrase, and accountability quietly disappeared inside the language. The insurance market noticed.

The renewal conversation nobody wants

Picture a mid-sized professional services firm. They deployed an AI customer service chatbot last year, the kind that handles routine questions and routes complex ones to humans. The chatbot, doing what chatbots sometimes do, made a commitment to a client that no human authorized. The client treated the commitment as binding, the firm declined to honor it, and the client filed suit. The firm called its insurance broker about the errors and omissions policy.

The broker explained the renewal that had quietly happened three months earlier. The new policy carried an endorsement excluding coverage for any claim arising from AI use. The endorsement was buried in the policy materials and the broker had flagged the renewal as routine. Coverage was denied, and the firm's exposure became the firm's problem.

This conversation is about to happen a lot. The conditions for it were set in motion by three separate insurance moves between June 2024 and January 2026.

W.R. Berkley went first. In June 2024 the company issued Form PC 51380, a flat AI exclusion attached to three of its management liability product lines: Directors and Officers coverage, Employment Practices Liability, and Fiduciary Liability. The exclusion is absolute. Any claim arising from AI use, by any person or entity, including third-party AI sitting inside software the insured never wrote, falls outside the policy. The form is widely cited in secondary commentary as covering D&O, E&O, and fiduciary liability, but the primary form caption lists EPL where commentary says E&O. That terminology slip in widely circulated legal analysis is itself a useful sign of how fast the literature is moving and how thin the verification chain has become.

Verisk went second. In January 2026 the company published ISO Form CG 40 47 01 26, an optional generative AI exclusion that carriers can add to commercial general liability renewals. Verisk's ISO forms underpin a large share of U.S. property and casualty policy language, so the form being optional matters less than the form existing. Once standard exclusion language is available, carrier adoption is a question of when, not whether.

AIG and Great American moved through state filings during 2025, seeking regulatory permission to issue their own AI-related exclusions. AIG stated publicly that it had no immediate plans to implement the exclusions it was filing for, which is consistent with carriers preserving the option to act later without restarting the regulatory clock.

The first market answer to AI exposure has been exclusion language rather than sophisticated governance pricing. That is what markets do when they cannot yet price a risk cleanly. They exclude first, watch the data come in, and price what they learn.

What the insurance scholarship is showing

Anat Lior published an empirical study of the AI insurance market in the Connecticut Insurance Law Journal in 2025 that traces this pattern in detail. Existing policy language often does not say anything explicit about AI, which leaves what the field calls silent AI coverage. Silent coverage is risky for both sides because the insured does not know whether a claim will be paid and the carrier does not know whether it will have to pay one. Both sides have an incentive to resolve the silence, and the resolution is happening now, sometimes through affirmative coverage language and sometimes through exclusions like Berkley's.

Two AI-specific products are already on the market. Munich Re sells a product called aiSure, which covers AI performance errors and certain legal and financial losses. Armilla operates as the world's first Lloyd's coverholder dedicated to AI insurance and covers AI underperformance, hallucination, and drift. Both products price the AI system itself or the AI's outputs. Neither product yet treats the human checkpoint as the central pricing variable. That gap is the argument this paper makes will close.

The legal scholarship has a name for what fills that gap badly. Crotoft, Kaminski, and Price wrote a paper in the 2023 Vanderbilt Law Review titled *Humans in the Loop*, and they called the bad version of human oversight a liability sponge. A liability sponge is a person stationed near an AI decision who lacks the time, the authority, the information, or the standing to actually decide. The AI runs at machine speed, the person clicks approve, and when something goes wrong the person absorbs the blame while the company absorbs the loss. Nothing about the system actually got governed. The paper's central argument is that putting a human near an AI decision is meaningless unless the human can really decide.

The pricing question is moving from model performance to governed consequence. The market is starting to ask who reviewed the output, what authority they held, what record they preserved, and whether the decision can survive an audit.

How risk becomes cost

The market has been through this loop before, in cyber.

Cyber insurance in the early 2010s was a mess of vague coverage and guesswork premiums. Then ransomware happened, claims got expensive, and underwriters started asking concrete questions about concrete controls. Did multi-factor authentication actually work in practice? Had the backup procedures been tested with restoration drills, or just configured on paper? Could the detection systems catch an intruder before the encryption finished running? Did the incident response plan exist beyond a binder on a shelf? Companies that could answer those questions in writing renewed at reasonable rates, while those that could not paid more, got narrower coverage, or got nothing.

The data from 2025 shows the cycle paying off. Coalition's 2025 Cyber Claims Report documented that Coalition policyholders had 73 percent fewer claims than the industry average benchmarked against NAIC-reported data. NetDiligence analyzed 10,402 claims from incidents between 2020 and 2024 and showed that companies with stronger preventative controls had measurably lower claim frequency and severity. AM Best reported that 2024 was the first year U.S. cyber direct premiums written declined since data collection started in 2015, a 2.3 percent drop driven in part by a market that had learned to price controls.

The same loop is now opening for AI, and the opening moves are what just happened with Berkley and Verisk. Silent coverage will give way to explicit wording, explicit wording will give way to dedicated AI insurance products, and underwriters will start asking specific questions about who reviewed what AI output and when. Companies with answers in writing will renew on reasonable terms, while the rest will not.

The six questions underwriters are starting to ask

Six variables determine where any AI deployment lands on the spectrum between excluded and insurable. Some of these questions are already being asked in current underwriting questionnaires. The rest are coming.

The first question is how the AI is being used. Does the system produce a final result that goes out into the world without a human in the loop? Or does it run to a human checkpoint that decides whether the output is acted on? The second configuration is far stronger for claim defense, because the organization can show that AI did not stand alone at the moment of consequence.

The second question is what being wrong actually costs. A wrong ad placement and a wrong medical recommendation have different cost profiles, and current AI insurance products do not yet have a standard instrument that prices the distinction with the granularity it deserves.

The third question is how the system handles its known failure modes. Every generative AI deployment carries some exposure to hallucination, drift, bias, and incomplete output. Underwriting will turn on whether the organization can show what it does to reduce those risks and how it records the reduction.

The fourth question is who is doing the verification at the human checkpoint. A campaign manager, a lawyer, a clinician, and a software engineer are not interchangeable as checkpoint authorities. The market does not yet have a standard instrument for measuring this, and it will need one.

The fifth question is whether the decisions are recorded in a way a third party can audit later. Does an audit trail exist that leads back to a named human at a discrete moment? Without that record, the firm has a weaker claim defense, a weaker coverage argument, and no clean answer when a D&O carrier asks what governance documentation existed before the deployment was authorized.

The sixth question is who is using AI that nobody approved. Gartner reported that 41 percent of employees acquired, modified, or created technology outside IT visibility in 2022 and projects that share will rise to 75 percent by 2027. The exposure inside that shadow IT footprint is real, and no standard AI insurance mechanism prices it cleanly yet.

Two pressure channels operating at once

The pressure on companies to answer these six questions is coming through two channels simultaneously, and they reinforce each other.

The first channel is corporate coverage at renewal. The Berkley exclusion and the Verisk optional endorsement give carriers tools to remove or narrow AI-related claims from legacy policy language. Companies that can document their governance architecture, their deployment mode, their checkpoint authorities, and their decision records will have a stronger underwriting position and a stronger claim defense. Companies that cannot will enter renewal with less to point to and more for the broker to explain away.

The second channel is personal exposure for the directors and officers who authorized the AI deployment in the first place. Derivative actions, regulatory inquiries, coverage disputes, and shareholder challenges will all ask the same question after an AI-related incident: what governance record existed before the decision to deploy was made? The simple version of that question is what changes behavior, because the answer has to be ready before the incident, not assembled after it.

Regulation is moving in the same direction. The NAIC adopted its AI Model Bulletin in December 2023 and more than twenty states have since adopted it or substantially similar guidance. The NAIC AI Systems Evaluation Tool has moved into multistate pilot use, which means state insurance regulators are turning AI governance into examinable evidence. The EU AI Act imposes obligations on deployers of high-risk AI systems for competent human oversight, six-month log retention, operational monitoring, and incident reporting. The EU Revised Product Liability Directive, adopted in late 2024 with a December 9, 2026 transposition deadline, brings software and AI into the modern strict liability frame for products and lowers the evidentiary barriers claimants face. Regulatory pressure and market pressure are now pointing the same way.

The bottom line

The three tiers the industry has been collapsing into one term are becoming three distinct insurance profiles. Ethical AI identifies the boundary and Responsible AI documents the system, but neither produces what an insurer or a court actually wants to see when a claim arrives. AI Governance produces that. It produces the named human, the defined checkpoint, the binding authority, and the audit record that together build a defensible claim narrative.

The market is starting to price the difference. The January 2026 exclusions are the opening move in a cycle that has already played out once in cyber and is now playing out again. Companies that act on the difference between ethics, responsibility, and governance now will renew at reasonable terms. The rest will spend 2026 having the renewal conversation nobody wants.

Acknowledgments

This working paper developed from a cross-disciplinary conversation hosted by The Intention Advantage, a working initiative of the Institute for Advertising Ethics (IAE), with the IEEE Standards Association serving as knowledge partner. The author thanks Andrew Susman, President of the Institute for Advertising Ethics, for the founding convening held April 22, 2026 at Frankfurt Kurnit Klein & Selz in New York, where the connections between advertising governance, agentic AI accountability, and the insurance pricing of human oversight were first surfaced for a cross-disciplinary audience. Contributions from IEEE Standards Association leadership, including the IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems, shaped the architecture of the argument developed here. Responsibility for the analysis, the framework, and any errors remains with the author.

Primary Sources

- Coalition Inc. (2025). 2025 Cyber Claims Report. Published May 7, 2025. <https://web.coalitioninc.com/download-2025-cyber-claims-report.html>
- NetDiligence. (2025). Cyber Claims Study 2025: A Data-Driven Analysis of 10,402 Cyber Insurance Claims from Incidents Occurring Between 2020 and 2024. Published September 17, 2025. <https://netdiligence.com/cyber-claims-study-2025-report/>
- AM Best. (2025). Best's Market Segment Report: 2024 Pricing Cuts in U.S. Cyber Generated First-Ever Reduction in Direct Premiums Written. Published June 23, 2025.
- NAIC. (2025). Report on the Cybersecurity Insurance Market. https://content.naic.org/sites/default/files/inline-files/2025_Cybersecurity_Insurance%20Report.pdf
- Lior, A. (2025). E/Insuring the AI Age: Empirical Insights into Artificial Intelligence Liability Policies. 31 Conn. Ins. L.J. 99. SSRN Abstract ID 5316376.
- Lior, A., & Madhok, S. (2025, December 9). Insuring the AI age. Willis Research Network Newsletter, WTW. <https://www.wtwco.com/en-nl/insights/2025/12/insuring-the-ai-age>
- Madhok, S., & Lior, A. (2026, February 27). AI liability in practice: What risk managers need to know now. Willis Research Network Newsletter, WTW. <https://www.wtwco.com/en-nl/insights/2026/02/ai-liability-in-practice-what-risk-managers-need-to-know-now>
- Crootof, R., Kaminski, M. E., & Price, W. N., II. (2023). Humans in the Loop. 76 Vand. L. Rev. 429.
- W.R. Berkley Corporation. Form PC 51380 00 (06-24) Artificial Intelligence Exclusion (Absolute). Primary form caption verified at Hunton Andrews Kurth document hosting. The form is marked as an amendment to base forms PC 51000 DO (02-18), PC 51000 EPL (02-18), and PC 51000 FLI (02-18), adding exclusion language to Section III of all Liability Coverage Parts. URL: <https://www.hunton.com/assets/htmldocuments/noindex/PC-51380-00-06-24-Artificial-Intelligence-Exclusion-Absolute.pdf>
- Anderson Kill, P.C. Insurance for AI Liabilities: An Evolving Landscape, as one representative of secondary legal commentary that describes the Berkley exclusion through D&O, E&O, and fiduciary liability terminology. The primary form does not list Errors and Omissions among the amended base forms.
- Verisk ISO Form CG 40 47 01 26 (effective January 1, 2026); EU AI Act Article 26 (deployer obligations); EU Revised Product Liability Directive (EU) 2024/2853 (transposition deadline December 9, 2026); NAIC AI Model Bulletin (adopted December 2023).

#Assisted using the HAIA Ecosystem

Term List

AI Governance A named human with binding authority at a defined decision checkpoint. Accountability leads back to a specific desk, a specific person, and a specific moment. Produces an auditable record. The only tier of the three that generates an insurable event.

Audit Trail A documented record of who reviewed an AI output, under what authority, at what moment, and what decision was made. The actuarial foundation of any AI insurance product. Currently absent from most enterprise AI deployments.

Checkpoint A defined stop in an AI workflow where a named human with domain competence reviews the output and makes a binding decision before the process continues. Distinct from a loop. A loop has no defined stop and no discrete record. A checkpoint does.

D&O (Directors and Officers) Liability Insurance covering the personal assets of corporate directors and officers against claims arising from their decisions. W.R. Berkley's absolute AI exclusion removed this protection for AI-related governance failures as of 2025. Directors who authorize AI deployments without documented governance checkpoints now face personal exposure in derivative actions.

Deployment Mode How an AI system is being used in practice. Automation producing a final result with no human checkpoint is one mode. Automation running to a named human who makes the consequential decision is another. The January 2026 exclusions distinguish between these modes for the first time in insurance history.

E&O (Errors and Omissions) Professional liability insurance covering claims arising from mistakes or failures in professional services. Now subject to generative AI exclusions under W.R. Berkley's absolute language and ISO Form CG 40 47 01 26.

Economic Override Pattern The structural tendency for deployment incentives to override governance regardless of institutional intent. Speed and cost efficiency consistently outweigh safety and accountability when the economic consequences of ungoverned deployment are not priced into the cost structure.

Ethical AI The normative foundation of the three-tier stack. Identifies harm vectors, establishes values, and draws lines where AI should not operate alone. Does not enforce anything. Does not produce a named human. Does not generate an insurable event.

Generative AI As defined in Verisk ISO Form CG 40 47 01 26: a machine-based system or model trained on data with the ability to create content or responses. Broad enough to capture chatbots, content generators, LLMs, and automated decision-making tools embedded in enterprise software.

Hallucination An AI output that is confidently stated but factually incorrect or fabricated. Present in every platform without exception. A primary driver of the January 2026 insurance exclusions. Not a reason to avoid AI. A reason to require a documented human review before a hallucinated output reaches a consequential decision point.

Human in the Loop (HITL) The practice of placing a human in an AI workflow. As currently used in insurance and governance literature, it describes presence without specifying authority, competence, or

record-keeping. Presence is not governance. A plumber in a medical AI review loop and a doctor in the same loop are not equivalent risk profiles. The market has not yet priced this distinction.

ISO Form CG 40 47 01 26 The Verisk Insurance Services Office optional endorsement, effective January 1, 2026, that allows US commercial general liability carriers to exclude bodily injury, property damage, and personal and advertising injury arising out of generative artificial intelligence. Underlies approximately 82 percent of U.S. property and casualty policies. The first standardized instrument to price ungoverned AI out of coverage.

Loss Data Historical claims records that actuaries use to set premiums. Currently absent for AI deployments because no standardized record-keeping requirement exists. The Lior (2025) study identifies this as the central obstacle to AI insurance market maturation. A governance audit trail generates loss data as a structural byproduct.

Multi-Factor Authentication (MFA) The cyber security control that became the reference precedent for how insurance markets drive governance adoption without legislation. 43 percent of organizations without MFA were denied cyber coverage entirely. The same mechanism is now activating for AI governance documentation.

Responsible AI The engineering of constraint. Testing, documentation, bias detection, hallucination monitoring, human in the loop as a process rather than a checkpoint. Produces traceability artifacts and proves a system behaves as intended. Does not produce a named human accountable for a specific decision. The ceiling of Responsible AI is the absence of individual human oversight.

Shadow AI The use of AI tools by employees that have not been approved or governed by the organization. Currently affects 41 percent of the workforce and is projected to reach 75 percent within three years. No current insurance policy prices this exposure. Three parties hold plausible liability with no clear coverage between them: the employee, the organization, and the AI provider.

Silent AI Coverage Traditional insurance policies that neither explicitly include nor exclude AI-related risk. Creates dangerous ambiguity about coverage. As documented in Lior (2025), the insurance market is resolving silent coverage into explicit exclusion rather than affirmative coverage, producing immediate gaps for organizations deploying AI without documented governance.

Three-Tier Stack The sequential architecture of Ethical AI, Responsible AI, and AI Governance. Not a menu from which organizations choose. A stack in which each tier provides a foundation the next tier builds on. The industry has collapsed all three into a single marketing term. The January 2026 exclusions began pricing the difference.